AAMIR ALI

LinkedIn | 7462-158-261 | RM8 3ED , London Uk | aamirali7596@gmail.com | GitHub

Skills

-**Programming Languages**: Python, SQL, Scala, C++
-**Databases**: CosmosDB, MSSQL, NoSQL, SnowSQL, kdb+
-**Cloud Platforms**: AWS (S3, EC2, EMR, Glue, Redshift), Snowflake, Azure (Databricks, Data Factory, Synapse), GCP, Terraform
-**Data Engineering Tools**: DBT, Apache Spark, PySpark, Spark Streaming, Kafka, Data Vault 2.0, CDC, ELT, Looker, Looker API, Apache Airflow (DAGs), Great Expectations, Azure DevOps
-**CI/CD**: Git, GitHub Actions, CI/CD pipelines, Azure DevOps, Agile/Scrum, TDD, Pytest, Automated Testing, L3 Support
-**Data techniques**: ETL, ELT, Data Modelling, Physical Modelling, Lakehouse Architecture, Schema Evolution, Data Lineage, Data Catalogue, Data Ingestion, Data Migration, Streaming, Data Warehousing, Data Analysis, CDC, Reconciliation, Semantic Layers, Rules Engines, GDPR

Experience

Senior Data Engineer **Build A Rocket Boy** *London, United Kingdom* **05/2025 - Current**

- Designed and developed scalable **PySpark** and **Apache Spark** data pipelines on **Databricks**, integrating 100+ source systems into a centralised **Snowflake** warehouse processing 120+ TB of structured and semi-structured data.
- Led **dbt** data model migration applying **TDD** practices and **CI/CD** pipelines via **Azure DevOps**, establishing target schemas and **automated testing** suites with **Pytest**, cutting deployment errors by 40%.
- Automated **ELT** workflows processing 5-8 TB daily from telemetry systems using **Python**, **PySpark**, and **Terraform** IaC on **AWS**, cutting ingestion latency by 35% and enabling self-service data discovery.
- Engineered Gen AI-powered analytics on **Databricks** with **Looker API** integrations and a centralised data catalogue, enabling natural language querying for 150+ business users and reducing ad-hoc query load by 60%.

Lead Data Engineer **Castleton Commodities Int.** *London, United Kingdom* **08/2024 – 04/2025**

- Architected scalable **PySpark** and **Kafka** streaming pipelines on **AWS**/**Azure** with **CDC**-based ingestion and **Apache Airflow DAG** orchestration, supporting trading desks processing 250+ TB of market data.
- Developed **Python** and **SQL**-driven reconciliation and **ELT** transformation pipelines for pre-trade market data (prices, positions, exposures, curves), supporting P&L reporting and VaR analysis.
- Built data quality validation frameworks and Power BI dashboards establishing full data lineage and a documentation library for all **Snowflake data warehouse** changes across 30+ front-office stakeholders.
- Provided **L3 support**, developed **Python** hotfixes for production **Databricks** incidents, and mentored junior developers on **PySpark** and **dbt** best practices, raising overall code quality standards.
  enabling informed decision-making for Analysts.
- Integrated REMIT/EMIR regulatory datasets into **Terraform**-managed governed data models, enforcing data privacy, access management, and **GDPR** compliance across all reporting and settlement layers.
-Integrated **regulatory reporting datasets (REMIT, EMIR)** into governed data models, ensuring full lineage, auditability, and compliance.

Lead Data Engineer **Sodexo** *London, United Kingdom* **05/2023 – 01/2024**

- Led **Databricks** and **PySpark**-driven **CDC** pipeline design for the AMETA region using **Azure** ADF, building scalable ingestion infrastructure that eliminated third-party tooling and reduced costs by 80%.
- Reported to CTO as engineering lead, building a data catalogue and technical documentation library for the fintech transaction platform, enabling data discovery across 1M+ daily user transactions.
- Designed **Agile/Scrum**-driven **ELT** pipelines via **Azure Data Factory** with **Apache Airflow DAG** orchestration, implementing schema evolution and partitioning strategies to optimise warehouse performance.
- Mentored junior developers on **Databricks** and **dbt** best practices; led **Data Vault 2.0** implementation with data quality rules engines, reducing data redundancy by 45% and accelerating delivery cycles.

Data Engineer **Wherescape USA** *(United States -Remote)* **07/2019 - 05/2023**

- Implemented **automated testing** suites using **Pytest** and **CI/CD** pipelines for BAT's **Snowflake** multi-cluster migration, delivering workload isolation, SLA compliance, and full **L3 support** coverage.
- Designed **Snowflake** external tables on **AWS** S3 for semi-structured data (JSON, Parquet) using **physical modelling**, partitioning, and schema evolution best practices, integrating IoT feeds via **ELT** pipelines.
- Migrated AMD Malaysia data warehouse from Informatica to Wherescape using **Agile/Scrum**, authoring **SQL** transformations and building incremental **ELT** pipelines loading 1B records with 80% efficiency gain.

Education

Bachelor of Computer Applications **Pune University** *Mumbai, India* **05/2016 - 04/2019**

•Bachelors in Computer Applications – 3.2 GPA