AAMIR ALI

LinkedIn | 7462-158-261 | RM8 3ED , London Uk | aamirali7596@gmail.com | GitHub

Skills

-**Programming Languages**: Python, Pandas, SQL, C++
-**Databases**: CosmosDB, MSSQL, NoSQL, SnowSQL, kdb+
-**Cloud Platforms**: AWS (S3, EC2, EMR, Glue, Redshift), Snowflake, Azure (Databricks, Data Factory, Synapse), GCP, Terraform
-**Data Engineering Tools**: DBT, PySpark, Kafka, Data Vault 2.0, CDC, ELT, Looker, LookML, Looker API, Apache Airflow (DAGs), n8n
-**CI/CD**: Git, CI/CD pipelines, Azure DevOps, Agile/Scrum, TDD, Automated Testing, L3 Support
-**Data techniques**: ETL, ELT, Data Modelling, Physical Modelling, Data Contracts, Data Lineage, Observability, Fault-tolerant Pipelines, Schema Evolution, Data Ingestion, Data Migration, Streaming, Data Warehousing, Data Analysis, CDC, Reconciliation, Semantic Layers, Rules Engines, GDPR

Experience

Senior Data Engineer **Build A Rocket Boy** *London, United Kingdom* **05/2025 - Current**

- Owned end-to-end data platform architecture, designing **dbt** data contracts and **semantic layer** models integrating 100+ source systems into a centralised **Snowflake** warehouse processing 120+ TB.
- Led platform migration to dbt Fusion, establishing modular **ELT** schemas, **CI/CD** pipelines, **automated testing** suites, and Architecture Decision Records (ADRs) to enforce engineering standards across the team.
- Engineered fault-tolerant **ELT** pipelines processing 5-8 TB daily using **Python**, **Terraform** IaC on **AWS**, with built-in observability and data quality monitoring, reducing pipeline failures by 35%.
- Designed ML-ready data foundations with **Looker API** and **LookML** semantic abstractions, enabling Gen AI-powered natural language querying for 150+ users with production-grade traceability and data lineage.

Lead Data Engineer **Castleton Commodities Int.** *London, United Kingdom* **08/2024 – 04/2025**

- Architected fault-tolerant **ELT** pipelines across **Snowflake**, **dbt**, and **AWS/Azure** with **CDC**-based ingestion and **Apache Airflow** DAG orchestration, supporting trading desks processing 250+ TB.
- Developed **Python** and **SQL**-driven reconciliation pipelines with clearly defined data contracts for pre-trade market data (prices, positions, curves), supporting P&L reporting and VaR analysis.
- Built rules-engine-driven data quality and observability frameworks with full data lineage tracking, delivering transparency into P&L and trade settlements for 30+ front-office and risk stakeholders.
- Owned production incident resolution as **L3** escalation lead, developed **Python** hotfixes, authored technical documentation and ADRs, and mentored engineers on system design and **SQL** best practices.
  enabling informed decision-making for Analysts.
- Designed enterprise-grade governance framework using **Terraform**-managed infrastructure, enforcing access control, data lineage, auditability, and GDPR compliance across REMIT/EMIR regulatory datasets.
-Integrated **regulatory reporting datasets (REMIT, EMIR)** into governed data models, ensuring full lineage, auditability, and compliance.

Lead Data Engineer **Sodexo** *London, United Kingdom* **05/2023 – 01/2024**

- Led end-to-end data platform design and **CDC**-based ingestion for AMETA using **Azure** ADF, **Databricks**, **PySpark**, eliminating third-party tooling and reducing infrastructure costs by 80%.
- Reported to CTO as senior technical lead, defining data contracts and architecture standards for a fintech transaction platform supporting reconciliation insights across 1M+ daily user decisions.
- Designed modular **ELT** pipelines via **Azure Data Factory** and **Apache Airflow DAG** orchestration with workflow automation, enabling fault-tolerant, schema-evolving data delivery to the warehouse.
- Defined technical roadmap for **Data Vault 2.0** **physical modelling **adoption; mentored engineers on system design and** SQL**, implementing rules engines that reduced data redundancy by 45%.

Data Engineer **Wherescape USA** *(United States -Remote)* **07/2019 - 05/2023**

- Delivered **automated testing** suites and **CI/CD** pipelines for BAT's **Snowflake** multi-cluster migration, owning production deployment with workload isolation ensuring SLA compliance and **L3 support**.
- Designed **Snowflake** external tables on **AWS** S3 for semi-structured data (JSON, Parquet) applying **physical modelling** and partitioning best practices, integrating IoT/supply-chain feeds via **ELT** pipelines.
- Migrated AMD Malaysia data warehouse from Informatica to Wherescape using **Agile/Scrum**, authoring **SQL** transformation scripts and building modular **ELT** pipelines loading 1B records with 80% efficiency gain.

Education

Bachelor of Computer Applications **Pune University** *Mumbai, India* **05/2016 - 04/2019**

•Bachelors in Computer Applications – 3.2 GPA